



# **EXERCISE 1**

# **HIGH PERFORMANCE COMPUTING**



Giulio Fantuzzi

# PROBLEM OVERVIEW AND METHODOLOGY




## PROJECT AIMS

- To assess the performance of various **openMPI** algorithms for specific collective operations
- To build a performance model for the latency




## COLLECTIVE MPI OPERATIONS

- Broadcast
- Barrier




## LOTS OF DEGREES OF FREEDOM

- Algorithm
- Message Size
- Number of MPI processes
- Allocation type
- Additional parameters (iterations, warm-up,...)




## ARCHITECTURE

- **2 THIN** nodes of the **ORFEO** cluster
- 2 sockets per node
- 12 cores per socket




## METHODOLOGY

- **OSU** benchmark to estimate the latency
- Bash scripts to automate data gathering phase
- **SLURM** workload manager
-  for data analysis and models

# BROADCAST ALGORITHMS


**1** LINEAR




**2** CHAIN




**3** BINARY TREE




# BCAST: ALLOCATION TYPE COMPARISON



(a) Linear algorithm




(b) Chain algorithm



(c) Binary Tree algorithm

- Socket and core allocations show latency “jumps” when changing node, while node allocation progresses seamlessly
- Similar considerations apply to socket changes
- Surprisingly, jumps occur at 16 and 32, hinting at additional factors influencing communication
- Lines convergence: Linear vs Chain algorithm
- Tree algorithm (expected) superior performance

# BCAST: ALGORITHMS COMPARISON



- With a small fixed size, algorithmic implementation has a minimal impact
- Tree algorithm and default configuration outperform linear and chain
  - Tree structure efficiently distributes the communication among its branches
  - Default configurations are robust and efficient in different environments, leveraging platform-specific optimizations
- Linear vs chain insightful points of comparison:
  - Chain superior performance, primarily attributable to its utilization of contiguous cores
  - Chain capacity to partition messages into chunks during transmission, a feature absent in the linear




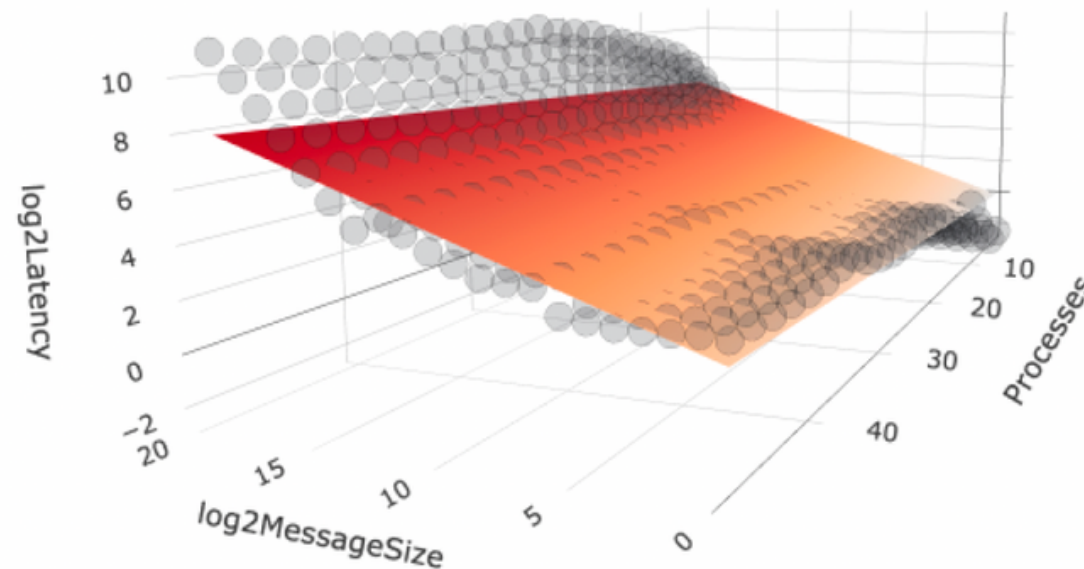

# BCAST: PERFORMANCE MODELS

- **INITIAL APPROACH:** estimate pt2pt latency and bandwidth to construct a model resembling the *Hockney's model*
- Unfortunately, results did not align well with the collected data, prompting me to explore statistical methodologies

$$\log_2(\text{Latency}) = \beta_1 \cdot \text{Number of Processes} + \beta_2 \cdot \log_2(\text{Message Size}) + \beta_3 \cdot \log_2(\text{Message Size})^2$$

Algorithm	$\beta_1$	$\beta_2$	$R^2_{adj}$
Linear	0.029815	0.321259	88,19 %
Chain	0.021957	0.306713	86,77 %
Binary Tree	0.016315	0.317449	86,58 %
Default	0.009586	0.321418	85,86 %

Algorithm	$\beta_1$	$\beta_2$	$\beta_3$	$R^2_{adj}$
Linear	0.0718655	-0.3022613	0.0355722	97,57 %
Chain	0.0646477	-0.3262914	0.0361133	98,00 %
Binary Tree	0.0600791	-0.3314744	0.0370215	98,38 %
Default	0.0525293	-0.3153369	0.0363273	97,74 %



# BARRIER ANALYSIS

- The analytical approach employed for Barrier closely follows the methodology followed for Broadcast
- Barrier aims for synchronization without involving message sizes —→ fewer degrees of freedom!
- My analysis includes the default MPI configuration, along with 3 additional algorithms:

- **LINEAR ALGORITHM**



All the processes report to a pre-selected root. Once every process has reported to the root, the root sends a releasing message to all participants

- **BRUCK ALGORITHM**

It requires  $\log_2(P)$  communication steps: at **step k**, process **r** receives a zero-byte message and sends it to process **(r-2<sup>k</sup>)** and **(r+2<sup>k</sup>)**, with *wrap around*

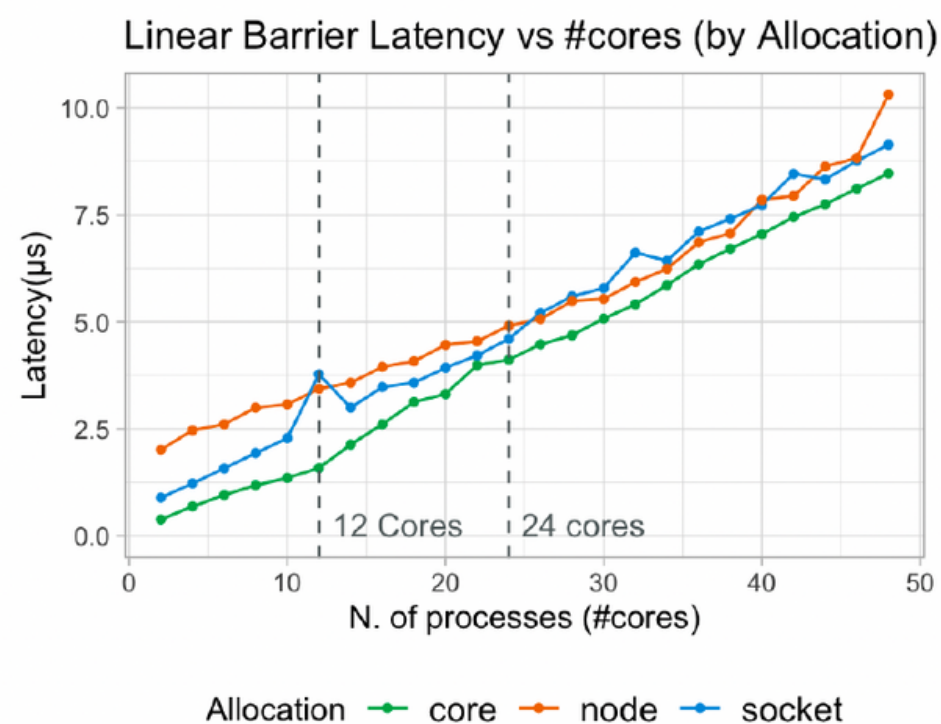
- **TREE ALGORITHM**

All the processes report to a pre-selected root. Once every process has reported to the root, the root sends a releasing message to all participants







# BARRIER: COMPARING ALLOCATION AND ALGORITHMS




(a) Linear algorithm




(b) Tree algorithm




(c) Bruck algorithm



(a) Configuration by core



(b) Configuration by socket



(c) Configuration by node



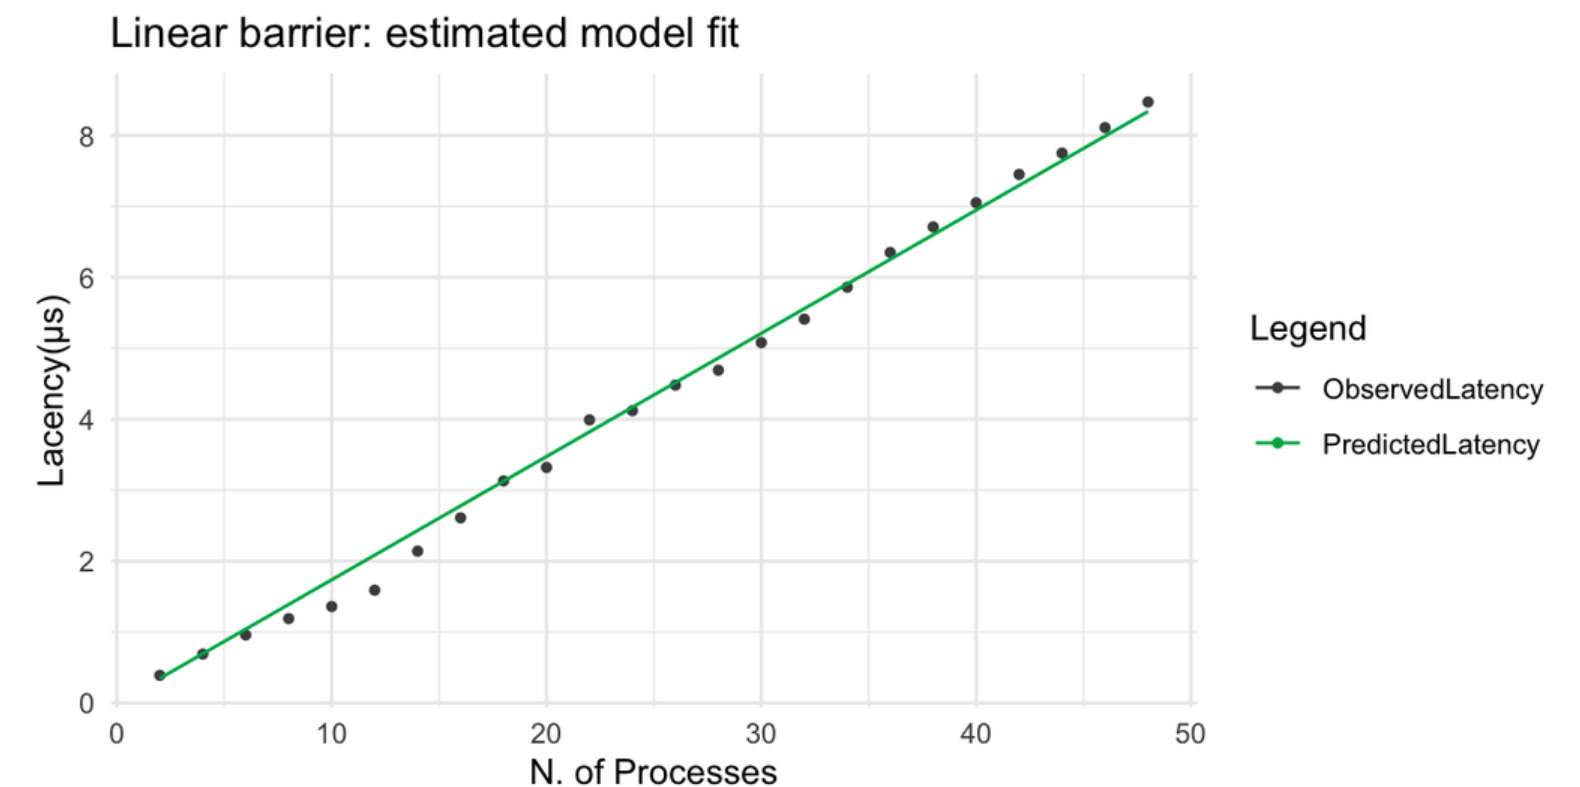
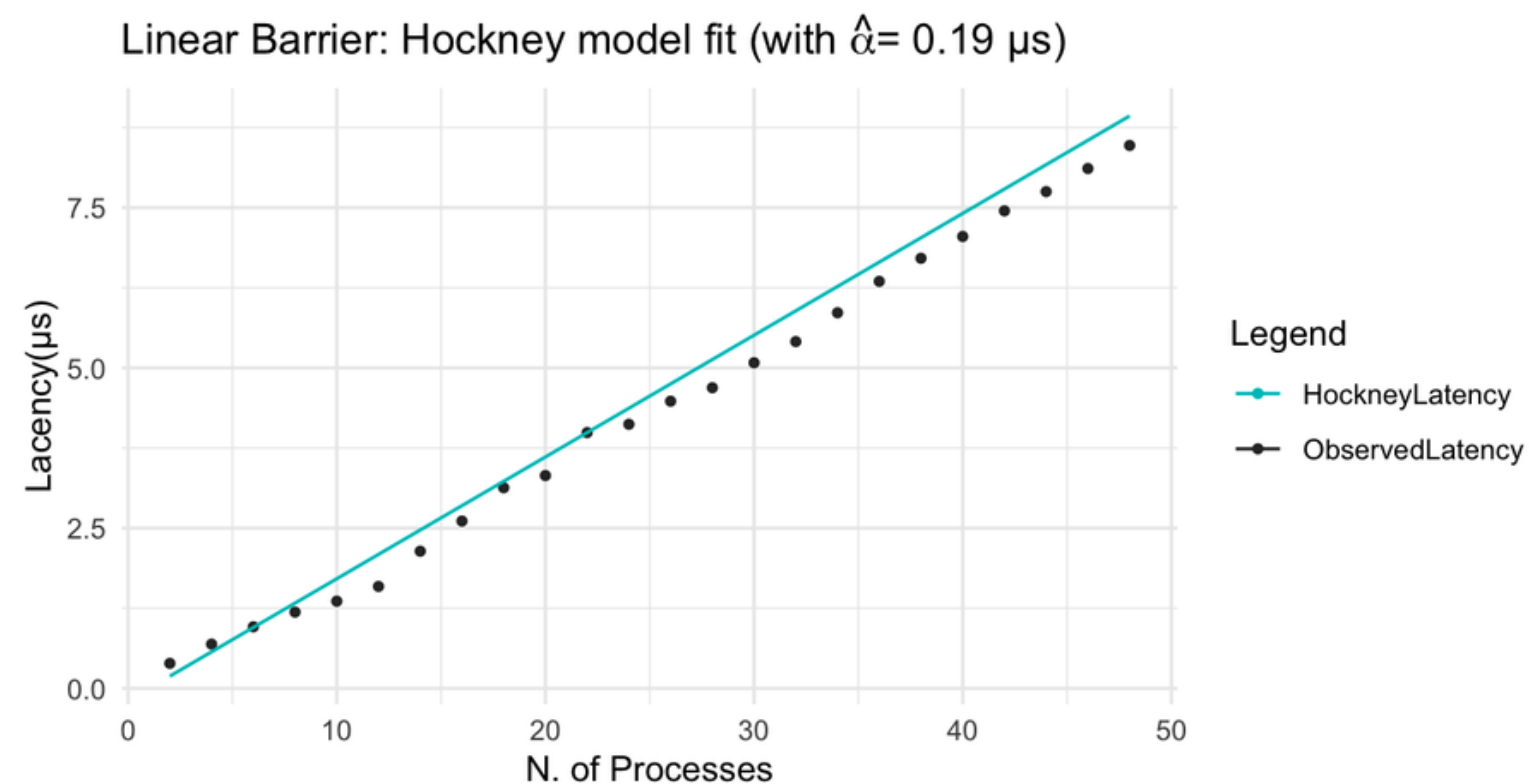
# BARRIER: PERFORMANCE MODEL FOR LINEAR ALG

- The chosen approach was to construct **algorithm-specific** models, in accordance with established practices in the literature
- For linear algorithm, I successfully integrated the estimated pt2pt latency (from **OSU**) into a Hockney model

$$\text{Latency} = \hat{\alpha}(P - 1) = \hat{\alpha}P - \hat{\alpha}$$

- As a curiosity, I implemented a linear regression model

$$\text{Latency} = \beta_1 \cdot \text{Number of Processes}$$



# BARRIER: PERFORMANCE MODEL FOR NON-LINEAR ALGS

- Attempts to extend the Hockney model to other barrier algorithms were unsuccessful
- For all the non-linear models, the introduction of a quadratic term improved the model's fit

$$\text{Latency} = \beta_1 \cdot \text{Number of Processes} + \beta_2 \cdot (\text{Number of Processes})^2$$

Algorithm	$\beta_1$	$\beta_2$	$R^2_{adj}$
Linear	0.1736878	/	99,86 %
Tree	0.1950677	-0.0019060	99,46 %
Bruck	0.1700697	-0.0012954	99,45 %
Default	0.1690773	-0.0010560	99,26 %

